

# Q-Learning on FrozenLake-v1: Effects of $\alpha$ , $\gamma$ , and $\epsilon$ on Learning

## 1. Environment and Setup

Environment: FrozenLake-v1 (4x4, is\_slippery=True).

States: 16 tiles (S, F, H, G). Actions: {Left, Down, Right, Up}.

Reward: 0 per step, 1 for reaching the goal, episode ends on hole or goal.

Episodes: 5000, averaged across 5 seeds.  $\epsilon$  decays from 1.0  $\rightarrow$  0.05 with rate 0.9995.

Exploration-exploitation controlled via  $\epsilon$ -greedy policy.

## 2. Algorithm: Q-Learning

Update rule:

$$Q(s,a) \leftarrow Q(s,a) + \alpha [ r + \gamma \max_{a'} Q(s',a') - Q(s,a) ]$$

Exploration:  $\epsilon$ -greedy, where random actions occur with probability  $\epsilon$ .

Learning parameters:  $\alpha \in \{0.1, 0.3, 0.5\}$ ,  $\gamma \in \{0.9, 0.99\}$ .

## 3. Results and Observations

Higher  $\alpha$  led to faster learning but less stable convergence.

$\gamma = 0.99$  outperformed  $\gamma = 0.9$ , showing that valuing future rewards helped in sparse reward settings.

Slower  $\epsilon$  decay improved long-term performance by maintaining exploration longer.

Below is a typical learning curve showing moving-average reward over episodes.

(Insert learning curve figure here: qlearning\_frozenlake\_curves.png)

## 4. Discussion

•  $\alpha$  (learning rate): Too low  $\rightarrow$  slow progress; too high  $\rightarrow$  oscillations.

•  $\gamma$  (discount factor): Higher  $\gamma$  improved stability and long-term goal pursuit.

•  $\epsilon$  (exploration): Gradual decay ensured sufficient exploration before exploitation.

Overall, balanced parameter choices ( $\alpha=0.3$ ,  $\gamma=0.99$ ,  $\epsilon_{\text{decay}}=0.9995$ ) yielded the best results.

## 5. Reflection

The experiment demonstrated the exploration-exploitation dilemma clearly. Excessive exploration delayed convergence, while premature exploitation trapped the agent in suboptimal policies. Future work could include comparing SARSA with Q-learning and testing on larger maps (e.g., 8x8 FrozenLake).

### References:

- OpenAI Gym Documentation: [https://www.gymlibrary.dev/environments/toy\\_text/frozen\\_lake/](https://www.gymlibrary.dev/environments/toy_text/frozen_lake/)

- Towards Data Science Q-Learning Tutorial:

<https://towardsdatascience.com/simple-reinforcement-learning-q-learning-fcddc4b6fe56>